Hate speech detection in Italian tweets

Group 4

1 INTRODUCTION

1.1 Motivation

Although hate speech is a phenomenon that existed long before the advent of the Web, communication via the Internet has allowed this phenomenon to spread and develop with peculiar characteristics tied to the nature of online media, such as anonymity on one side, and velocity and breadth of proliferation on the other. Given the relevance and dangerous potential of this social phenomenon, there is widespread interest in recognizing and detecting hate speech online, also on an institutional level.

In May 2016 the European Commission, citing the chilling effect of hate speech on the democratic discourse on online platforms, agreed with Facebook, Microsoft, Twitter and YouTube on a "Code of conduct on countering illegal hate speech online" [3]. In this document the major IT Companies made a public commitment to:

1. Have in place clear and effective processes to review notifications regarding hate speech on their platforms;

2. review the majority of valid notifications of hate speech in less than 24 hours.

The most recent evaluation campaign of the Code of Conduct [7] assessed in general a worse performance of the IT Companies especially in this last respect (only 64.4% of notifications were reviewed in less than 24 hours), confirming a troubling downward trend already observed in 2021.

Accurate, precise and efficient automated systems of hate speech detection are necessary for combating the phenomenon both with proactive action and in the process of reviewing user notification of occurrences of such violations.

Moreover, such systems should prove valuable for providing the data that is still needed to investigate the ecosystem of hate speech online and the magnitude of the phenomenon.

1.2 Project goal

The goal of this project is to identify characteristics in the text which allow for the detection of Italian tweets containing hate speech (this category is meant to include expressions of racism, xenophobia and terrorist propaganda). So, we can define this task as a binary classification.

	HS	Not HS	Total
Train	2766	4073	6839
Test Tweets	622	641	1263
Test News	181	319	500

Table 1: Distribution of Hate Speech labels.

A secondary goal is to evaluate how insights obtained by our model can generalize across textual domains. In order to do that, we tested our models, trained exclusively on tweets, against Italian newspaper headlines.

The project is based on the main task of the HaSpeeDe 2 shared task presented at EVALITA2020 [8].

2 METHODS

2.1 Data

We used the datasets made available[1] by the organizers of the Evalita 2020 shared task HaSpeeDe 2 [8].

The train set consists in 6839 Italian Tweets posted between October 2016 and May 2019 annotated for presence of hate speech. The test set includes both a corpus of 1263 Italian Tweets posted between January and May 2019 and a corpus of 500 Italian newspapers' headlines retrieved between October 2017 and February 2018 from the online editions of *La Stampa*, *La Repubblica*, *Il Giornale* and *Liberoquotidiano*. This second corpus allowed us to appraise how well the application generalizes across textual domains.

Table 1 shows the distribution of hate speech labels in the training set and in each of the test sets.

The data sets are available in TSV format and contain three features for each record:

- `id`: numeric identifier for each document

- `text`: the body of the document

- `hs`: boolean value, whether the document contains HS (1) or not (0).

- `stereotype`: boolean value, whether the document contains a stereotype (1) or not (0).

As part of the preprocessing performed by the organizers of the task, mentions and URLs recurring in the original documents were replaced with `@user` and `URL` placeholders.

---

1 https://github.com/msang/haspeede/tree/master/2020

In order to model the semantic properties of the documents, we used the Italian Twitter embeddings [1] lexicon computed by ItaliaNLP Lab on a corpus of 50,000,000 tweets using the word2vec[2] toolkit. This lexicon consists of embeddings with 128 dimensions for 1,188,949 tokens that were computed with the CBOW model with a symmetric context window of 5 tokens.

The embeddings were made available[3] by the ItaliaNLP Lab as a SQLite database containing a single table `store` with 130 columns:

- `key`, representing the token;
- one column for each dimension of the embedding;
- `ranking` storing the frequency rank of the token.

Finally, we used two resources for identifying offensive words: a lexicon of 500 bad words for the Italian language made available on GitHub[4] and the Italian section of the Revised HurtLex lexicon [10] consisting of 7920 words. In the Revised HurtLex lexicon, each headword is annotated with an offensiveness level score derived from ratings provided by a large number of annotators. This resource was made available by the authors in TSV format[5] and was revised by us in order to solve some data quality issues.

2.2 Automatic text annotation

We automatically annotated the documents using a *Stanza* [6] pipeline: this operation consisted of:

- Tokenization
- Part of speech tagging
- Lemma extraction
- Dependency parsing

Before this step, we preprocessed the text in order to remove mentions, urls, and normalize punctuation: in particular, sequences of multiple full stops were reduced to one full stop, since they were creating issues with the sentence tokenization. As for hashtags, since often single words hashtags are integrated in the syntax of the sentence, we decided to remove only the hash mark and preserve the textual part of the hashtags.

Our choice of using *Stanza* was motivated by the possibility of using a combined model based on several different treebanks for the Italian language. This allowed us to get a better coverage of the language, which is particularly important when working on a non-standard variety such as tweets or, in general, computer-mediated writing. In particular, the combined model is based on four treebanks: ISDT, VIT, PoSTWITA and TWITTIRO. The latter two were developed specifically for the annotation of tweets.

2.3 Linguistic profiling

With the objective of identifying, if possible, stylistic features useful for hate speech detection, we developed 91 different features for the linguistic profiling of the text, belonging to five different categories:

- **Raw text properties:** text length in terms of number of tokens and number of sentences. Average length of sentences and words as a really simple proxy of syntactic complexity;

- **Social media language properties:** distribution of URLs, hashtags, mentions and all uppercase words;

- **Lexical variety** expressed in terms of Type/Token Ratio;

- **Morpho-syntactic properties**, namely the distribution of all grammatical categories and lexical density (proportion of content words in the text);

- **Syntactic properties:** distribution of all dependency relationships and relative order of core arguments and verb in sentences. Another aspect that received some attention was the distribution of nominal sentences, since the literature suggested a strong association between slogan-like nominal utterances and hateful content [2].

In order to get some insights into the challenges associated with the cross-domain task, we scraped headlines from the online editions of *La Stampa*, *La Repubblica*, *Il Giornale* and *Liberoquotidiano*. We filtered the headlines for the presence of the same small set of neutral keywords associated with immigrants, muslims and Roma used for the creation of the original dataset of tweets [4].

We thus obtained 748 newspapers headlines, which we passed to the automatic text annotation pipeline and analyzed stylistically. We then trained a Linear Support Vector Machine on the linguistic profiling features in order to determine which features contributed the most in distinguishing between tweets and headlines. The classifier was able to distinguish tweets from headlines with near certainty, obtaining a macro-F1 of 0.949 in 5-fold cross validation.

As expected, social media properties like the number of mentions, hashtags, URLs and the use of uppercase, were mainly responsible for the success of the classifier. In addition to that, the feature importance analysis seems to indicate that tweets tend to be longer and have longer sentences than headlines; on the other hand, headlines seem to use more punctuation and more proper nouns, as well as exhibiting higher lexical variety (but since the Type/Token Ratio measure is highly sensitive to text length, this is probably a spurious result).

Figure 1 shows the distributions of several linguistic profiling features (document length in tokens, average sentence length, distribution of punctuation, distribution of proper nouns and Type/Token Ratio by form)
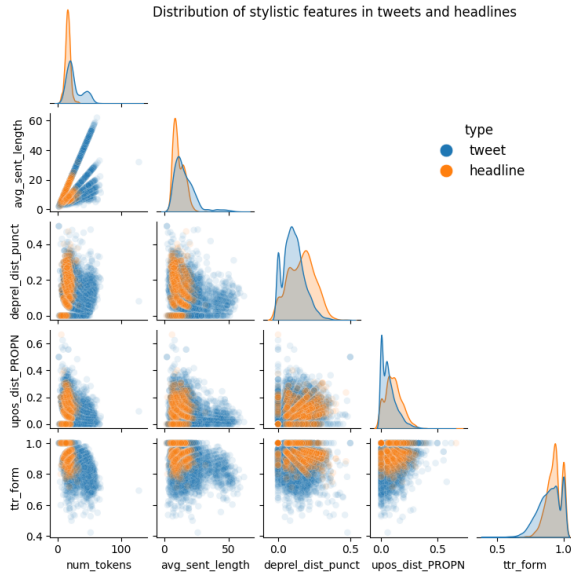
Distribution of stylistic features in tweets and headlines

Figure 1: Comparative stylistic analysis of tweets and scraped headlines.

for tweets and headlines. Although the distributions are overlapping, it is possible to discern clear trends that confirm the results of the feature importance analysis.

As we have seen, tweets and newspapers' headlines differ significantly in terms of stylistic features. Therefore, stylistic features that are effective for detecting hate speech in tweets are unlikely to yield good results when applied to such a different textual domain.

The lemmatized text was also used for the identification of offensive words from the HurtLex lexicon in the documents: we computed a global offensiveness score for each document adding together the scores of all offensive words recurring in the text. We then normalized this score by text length in order to avoid overestimating the offensiveness of longer documents. The necessity of this ulterior step was made clear when we noticed that the scraped newspapers' headlines tended to be shorter and yield lower non-normalized offensiveness scores than tweets.

2.4 Ngrams extraction and modelling

In this section, we describe the process used to extract n-grams from tweets to exploit them as features.

The process begins with a pre-processing step applied to the raw training dataset, implemented through the `clean_df function`, which consists of three inner functions. The first inner function, `count_uppercase`, calculates the number of uppercase characters in the dataset. When we considered the stylistic features, we successfully replaced this simple counter with a function that computes the ratio between uppercase and lowercase characters.

Additionally, we defined a custom punctuation set that extends the standard `string.punctuation` by including special characters encountered in tweets. After applying this function, the dataset includes the orig-

inal columns, text and hs, along with a new column, `uppercase_count`.

The second inner function, `split_hashtags`, utilizes the *WordSegment* library to split words within hashtags. This function is stochastic, meaning it may split hashtags differently during each execution. Based on our observations, the function performs well in identifying and splitting longer words within hashtags but struggles with stopwords and shorter words. This second inner function is incorporated into the `process_hashtags` function, which tokenizes hashtags when more than one word is present. An additional condition was added to handle words longer than 10 characters, although this could occasionally split valid long words incorrectly. However, given the rarity of such words in the dataset, this approach is effective for identifying incorrectly joined words. At this stage, the dataset includes a new column, `text_processed`.

Subsequently, all text was converted to lowercase, mentions, URLs, and numbers were removed, and the tweets were tokenized using *Stanza* [6]. Tokens were then further processed by removing stopwords and punctuation. After tokenization, two additional columns were added: one reporting the final number of tokens per tweet and another counting the occurrence of bad words, if any, from the list of 500 Italian badwords.

Once the tokenized version of the tweets was prepared, we lemmatized the tokens to ensure higher consistency and reduce the number of final n-grams. Given the brevity of tweets, we limited the n-grams to unigrams and bigrams.

The final dataset consisted of 6837 rows and 6407 features, including the original tweet, tokenized version, lemmatized version, the number of uppercase characters, the count of bad words, and extracted unigrams and bigrams. Before extracting additional features and embeddings, we conducted initial experiments by training models directly on this dataset. Several models were tested, including Random Forest, SVM, XGBoost, and Logistic Regression using as features only n-grams and badword counts. After applying Random Search for hyperparameter tuning, the best performance was achieved with Logistic Regression, yielding a macro average F1 score of 0.743. In this step, we did not use SelectKBest to select the top k features, as this would have prevented the model from being properly applied to the test data. Instead, we used all the features.

The same preprocessing steps used for the training set were applied to both test sets: one containing tweets and the other containing 500 headlines from Italian newspapers. To ensure consistency between the test datasets and the training set, we used the dictionary, the Bigram model, and the TF-IDF model fitted on the training set. As a result, the final test datasets have the same number of features as the training set, allowing the model trained on the training dataset to be correctly applied to the test data. Specifically, if an n-gram appears in both the training and test sets, the corresponding column will report non-zero values; otherwise, the column will consist entirely of zeros.

Moreover, using the same TF-IDF model for both training and testing ensures that the weight of each n-gram remains consistent.

2.5 Handling Emoji

Analyzing emoji in Italian tweets is essential for understanding expressiveness, sentiments, and the main topics shared online. Emoji represent a unique form of communication that complements textual language with visual symbols to convey emotions, emphasis, and cultural context. Our objectives are therefore to:

- Identify emoji usage patterns in Italian tweets.

- Analyze the frequency and distribution of emoji.

- Explore the semantic meaning of emoji in relation to textual content.

- Apply clustering techniques to categorize tweets based on emoji usage.

2.5.1 Dataset and Preprocessing

To ensure the quality of the analysis, the dataset was pre-processed using the following steps:

1. **Removal of superfluous elements**: Mentions (@user), URLs, hashtags, and non-decodable characters were removed.

2. **Text normalization**: All text was converted to lowercase, and redundant spaces were removed.

3. **Tokenization**: The text was separated into individual words or symbols for detailed analysis.

4. **Emoji extraction**: The emoji library was used to identify and isolate emoji symbols from the text.

2.5.2 Extraction and Analysis

We performed two different extractions — one for emoticons and another for emoji. The former required designing a pattern to identify them, while the latter was handled through tokenization and recognition patterns provided by the emoji library. The following analyses were conducted:

- **Emoji counting**: Calculation of the frequency of each emoji in the dataset.

- **Meaning mapping**: Linking emoji to their descriptive names and semantic meanings.

- **Emoji distribution**: Statistical analysis to identify the most common emoji.

- **Relationship with textual content**: Exploration of the correlation between emoji usage and the topics discussed in tweets.


Figure 2: Distribution of the most frequent emoji in tweets.

2.5.3 Visualization and Results

The main visualizations developed include:

1. **Emoji frequency**: A bar chart showing the most used emoji in the dataset.

2. **Clustering**: Use of algorithms such as K-Means and PCA to identify groups of tweets based on common emoji usage patterns.

As we can see from the bar chart, the most frequently used emoji is the enraged face. This visualization provides valuable insight into the trends among the tweets under analysis. As expected, the most prominent sentiment is anger. Subsequent results are also significant, showing the laughing face and then the "IT" pair of emoji. When combined, these emoji might suggest the prevalence of sarcasm, anger, and blame directed toward the government.

2.5.4 TF-IDF for Emoji and Emoticon Analysis

One key technique was the use of TF-IDF (Term Frequency–Inverse Document Frequency) to analyze the meaning and importance of both emoji and emoticons in relation to the textual content of the tweets. The process was structured as follows:

1. **Text tokenization**: Each tweet was divided into tokens, including both emoji and emoticons.

2. **Generation of n-grams**: For each emoji-emoticon, n-grams (from 1-gram to 5-gram) were generated to analyze combined usage patterns.

3. **Frequency calculation**: The relative frequency (TF) of each n-gram was calculated within each tweet.

4. **Inverse weighting calculation**: The relative importance of each n-gram in the corpus was calculated, penalizing those present in all tweets (IDF).

5. **Vectorization**: Each tweet was represented as a TF-IDF vector, with each dimension corresponding to an n-gram.

6. **Parameter optimization**: The optimal value of n (the length of the n-grams) was chosen based on total frequency.

Implementing TF-IDF made it possible to distinguish generic symbols with common meanings from those that are more specific and reflective of particular themes or emotions.

2.5.5 Results of the TF-IDF Analysis

Applying TF-IDF highlighted the following main n-grams with their corresponding TF-IDF values (cf. Table 2).

N-gram	TF-IDF
enraged_face	0.003181
italy	0.002182
face_with_tears_of_joy	0.002172
thinking_face	0.001885
xo	0.001316
pile_of_poo	0.001268
rolling_on_the_floor_laughing	0.001136
beaming_face_with_smiling_eyes	0.001036
d	0.001024
enraged_face_enraged_face	0.000979

Table 2: Top n-grams of emoji names with corresponding TF-IDF values.

### 2.5.6 Comparison of Emoji in Tweets With and Without Hate Speech

A comparison of the emoji used in tweets with and without hate speech was carried out to analyze differences and specific patterns. The relative frequencies of emoji were calculated in both datasets:

- **Tweets with Hate Speech (HS)**: Includes tweets where hs = 1.

- **Tweets without Hate Speech (NHS)**: Includes tweets where hs = 0.

Emoji frequencies were represented graphically for each group (cf. Figure 3).

As we can see in Figure 3, emoji such as "Italy," "face_with_tears_of_joy," "enraged_face," etc., appear in both rankings and also at the top positions. This suggests that the enraged face emoji can sometimes be used to express emotions of disappointment, political fatigue, and distrust regarding the topic of immigration. Conversely, the presence of the laughing emoji in the first ranking is indicative of malevolent sarcasm, mockery, and frustration. A similar observation applies to the Italian flag emoji, which is often used in both types of tweets.


Figure 3: Top 10 Emoji in Tweets With and Without Hate Speech.

Figure 4: Elbow Method for Determining the Optimal Number of Clusters.

Figure 5: K-Means Emoji Clustering (k=4) – 3D Visualization.

2.5.7 Tweet Clustering

Clustering was performed to categorize tweets based on emoji and emoticons. The main steps were:

1. **Data transformation**: We used the TF-IDF vectors based on the emoji names created before.

2. **Determining the optimal number of clusters**: The *Elbow* method was used to find the optimal value of k. Figure 4 shows the Elbow method plot.

3. **Dimensionality reduction**: `Principal Component Analysis` (PCA) was applied to reduce the data to three dimensions for visualization.

4. **Clustering**: The `K-Means` algorithm was used to identify four distinct clusters based on emoji usage.

Figure 5 shows the cluster distribution in a three-dimensional space based on the computed principal components.

The results show that emoji are not easily clusterable. Even though the optimal k found is 4, we end up with only 3 actual clusters. This is due to the nature of the available data. The division seems to be on the vertical layer and, based on the trend of the obtained data, we can hypothesize that the blue cluster contains emoji

	Features	Macro F1
1	Avg over all tokens	0.7443
2	Avg over all N, V and Adj	0.7255
3	Concat. of mean vects (N, V, Adj)	0.7175

Table 3: SVM classifiers based on different word embeddings aggregations (5-fold cross-validation).

with negative sentiment, the yellow cluster contains those with positive sentiment, and the magenta cluster comprises neutral emoji used in tweets featuring hate speech. These are merely hypotheses, since we are not able to derive more relevant information.

2.6 Embeddings

Among the approaches we decided to implement, we opted for an SVM classifier trained using embeddings-based textual representations.

We first tested three different representations based on the same lexicon of precomputed static word embeddings, the Italian Twitter Embeddings [1] by the ItaliaNLP Lab, differing by method of aggregation.

In order to do so, we loaded the embeddings and our preprocessed and annotated training documents, which we further normalized using the functions provided by the ItaliaNLP lab. We could then build the vocabulary of the training set using the embeddings. Given the characteristics of this type of non-standard language, in particular the high level of both interpersonal and intrapersonal variation, we expected some out-of-vocabulary words, which ended up being 2703 out of 22709, meaning that almost 12% of words did not have an embedding representation. These words are mostly words that haven't been correctly tokenized (especially multi-words hashtags), sequences of emojis, and typos. Had we used non-specialized embeddings, we assume that this coverage would have been significantly lower, posing an important issue for the model's training.

Using a 5-fold cross-validation process and comparing the resulting macro-averaged F1 scores, we determined the best performing model to be the one representing each tweet with the average of all its word embeddings. It performed better than using the aggregation of only lexically-full tokens (nouns, verbs and adjectives), both by mean and by concatenation of the mean vectors for the three aforementioned grammatical categories (cf. Table 3).

We then followed the same methodology to test representations based on contextual embeddings extracted from two different Italian pre-trained BERT models. [6]

1. **dbmdz/bert-base-italian-cased** or (**Italian BERT**) [9]: trained on a corpus consisting of a Wikipedia dump and texts from the OPUS corpora collection, for a total of 2,050,057,573 tokens.

---

6 **BERT** (Bidirectional Encoder Representations from Transformers) is a bidirectional Large Language Model (LLM) based on an encoder-only transformer architecture. This Deep Learning model was presented by Google in 2018.

	[CLS]	MEAN POOLING
Italian BERT	0.6814	0.7242
AlBERTo	0.7439	0.7465

Table 4: Mean macro F1 score of SVM classifiers based on different contextual embeddings' poolings (5-fold cross-validation).

2. **m_polignano_uniba/bert_uncased_L-12_H-768_A-12_italian_alb3rt0** (or **AlBERTo**) [5]: the first Italian BERT model for Twitter language understanding, trained on Italian tweets.

For both models, we tested two methods of *pooling* the tokens' embedding representations generated by the model after processing the input sequence into a single sentence embedding for each tweet (or, more generally, document).

1. **[CLS] pooling**: the model's tokenizer adds the special token [CLS] at the start of every sequence. During the training process of the Large Language Model, its embedding (or hidden state) will capture sequence-level information. By extracting the [CLS] token's embedding of each document from the final hidden layer, we obtain a representation that can serve as a sentence-level embedding.

2. **Mean Pooling with Sentence Transformers**: we used a Sentence Transformer, a model that generates a single fixed-length vector representation for an entire sequence. It does so via *mean pooling*, i.e. by computing the mean of all tokens' embeddings in the last hidden layer.

The Italian BERT yielded worse performances (cf. Table 4), which can be explained by the fact that it is a general model that has not been trained for this specific domain. AlBERTo, on the other hand, provided better results, with the SVM model based on its mean pooled embeddings having the highest mean macro F1 score (0.7465).

2.7 Models

The train set used for this phase of exploring various classifiers consists of four main components: unigrams and bigrams, embeddings, stylistic features, and the offensiveness score. Depending on which embeddings dataset is used, the dimensionality may change, but it generally consists of approximately 7,000 total features.

After tuning the hyperparameters of several classifiers (**Logistic Regression**, **SVM**, **Random Forest**, **XGBoost**, **Light-GBM**), **Light-GBM** emerged as the best performer, achieving an average F1-macro of 0.774 on 5 folds of cross-validation. Moreover, it proved to be the fastest. Thus, we employed this first version of the model to carry out several feature-focused analyses.

For instance, with respect to feature importance: although any of the four datasets manages, taken individually for training, to yield an average F1-macro of at least 0.75, the classification model primarily relies on


Figure 6: Performance of LGBM classifier vs number of features used by the classifier

the embeddings when all datasets are combined. It also considers a few other interesting features, such as the unigrams 'rom' and 'terrorism', the stylistic features 'uppercase_words_dist' and 'lexical_density', as well as 'offensiveness_score' and 'badword'.

Moreover, the calculation of feature importance revealed that the model does not benefit at all from many of our features, as it achieves its highest result (average F1-macro of 0.784) using only the first 100, as shown by the plot in Figure 6.

Since embeddings appeared to be the most useful feature set, we trained the model on all the embeddings datasets we had computed to identify any noticeable difference. It resulted that the best performing embeddings are the mean pooled embeddings extracted from AlBERTo using Sentence Transformers.

With regard to the offensiveness score, we found that the score normalized by the length of the tweet yielded better results than both than the non-normalized score and the combination of the two scores.

Lastly, using the insights we had gained, we conducted a more comprehensive randomized search to retune the Light-GBM model's hyperparameters in order to have a definitive model to use for testing. The model reached an average F1-macro of 0.778 on the validation set. For the cross-domain task, given the significant stylistics differences between tweets and newspaper headlines, we decided to perform a second tuning to train another model excluding stylistic features.

2.8 BERT

In this section, we describe the implementation of the same two Italian pre-trained BERT models used in 2.6 to extract contextual embeddings, Italian BERT and AlBERTo. The aim is to verify if the implementation of these models as binary classifiers, opportunely fine-tuned for the task at hand, proves to yield better performances than the more traditional models implemented earlier.

After loading the datasets, we loaded the pre-trained model's tokenizer using the Hugging Face Transformer library. We then prepared the text by adding the [CLS] token at the beginning of each sequence, to allow for correct processing by the tokenizer, which converts

Figure 7: AlBERTo Training and Validation Loss trend (LR = 2e-5, Epochs = 4).

the sequences in tokens that match its training data. Since BERT exploits fixed-lenght encoding, we set the maximum sequence lenghts for our sentences (128 for the training set, 512 for the test sets). We also added the required [SEP] token at the end of each sequence. The BERT tokenizer was then used to convert each token into an integer index that points to the proper embedding in the BERT vocabulary. Sentences shorter than the fixed maximum lenght were padded with zeros. Each token also received an attention mask, with 1 if it is a real token and 0 if it is padding.

After reserving 10% of the training set for validation, we created PyTorch Dataloaders for each set (train, validation, and the two test sets), to provide batches of 32 data sequences for the model to process in parallel.

We then proceeded to load the pre-trained model with a single linear classification layer added on top. We retrieved the model's training parameters and set the remaining hyperparameters for tuning as shown in Table 5, including the weight decay and learning rate for the Adam optimizer employed. The model was then trained and evaluated after each epoch.

For both models, we first tested 4 epochs with different learning rates for the Adam optimizer. However, plotting the loss trend on the training and evaluation dataset per epoch revealed that, while the training loss consistently decreased, the validation loss would increase after the second epoch (cf. an example in Figure 7) . This is a signal of overfitting: the model is memorizing the training data rather than learning from it, which hurts its capability to generalize on the validation set. Therefore, we decided to set the epochs to two, even if the training loss was still on a decreasing trend.

Parameter	Values
Epochs	2
Learning rate	2e-5
Weight Decay	0.01
Batch-size	32

Table 5: BERT's hyperparameters for tuning

3 RESULTS

Table 6 reports the performance of the tested models in terms of Macro-F1, comparing them to two baselines provided by the organizers of the shared task:

- A classifier returning the most frequent class (Baseline_MFC);

- A Linear SVM using TF-IDF of unigrams and 2-5 char-grams (Baseline_SVC).

Model	Twitter	News
Italian BERT fine-tuned	0.7696	0.6706
AlBERTo fine-tuned	0.7616	0.6838
LGB	0.7463	0.6656
LGB w/o style	0.7298	0.6709
SVC AlBERTo embeddings	0.7242	0.6674
Baseline_SVC	0.7212	0.621
Logistic Regression Ngrams	0.4004	0.5093
Baseline_MFC	0.3366	0.3894

Table 6: Macro-F1 of models tested on both test sets.

3.1 Logistic Regression with TF–IDF of Ngrams

Our Logistic Regression achieved a macro-average F1 score of 0.4004 on the tweets dataset and 0.5093 on the headlines dataset. In both cases, the low F1 score can be attributed to the low recall for class 1.

The results obtained from only the n-grams were poor. The seemingly better performance on the headlines dataset can likely be attributed to its more balanced class distribution.

3.2 Linear SVC with AlBERTo sentence embeddings

Detailed results for both test sets are reported in tables 7 and 8.

The macro F1 drops from 0.7465 observed in validation to 0.7242 for the in-domain task, but all metrics are balanced between the two classes. The model seems to be able to generalize well on tweets and has a stable performance.

As expected, the performance drop is more significant for the out-of-domain task, with a macro F1 of 0.6674. The model performs worse for class 1, especially with respect to recall, which means that it struggles with detecting hate speech in news headlines.

	precision	recall	f1-score
0	0.7513	0.6833	0.7157
1	0.7015	0.7669	0.7327
accuracy			0.7245
macro avg F1			0.7242

Table 7: Classification Report on Tweets test set, main metrics.

	precision	recall	f1-score
0	0.7417	0.8370	0.7865
1	0.6286	0.4862	0.5483
accuracy			0.7100
macro avg F1			0.6674

Table 8: Classification Report on News headlines test set, main metrics.

3.3 Light–GBM

Except for the fine-tuning of the two BERT models, Light-GBM obtained the best performance on both test sets, clearing both baselines. The model trained on all features obtained a macro F1 of 0.7463 for the in-domain task and of 0.6656 on the cross-domain task. As expected, excluding the linguistic profiling features yielded a better performance for the cross-domain task (0.6709), but lowered significantly the macro F1 in the in-domain task (0.7298).

3.4 BERT

To fine-tune our two pre-trained BERT models, Italian BERT and AlBERTo, we carried out the steps described in 2.8. We then proceeded to test both models against our two test sets. The main metrics from the final classification reports are presented in Tables 9,10,11 and 12.

	Precision	Recall	F1-Score
0	0.7796	0.7613	0.7703
1	0.7598	0.7781	0.7689
accuracy			0.7696
macro Avg F1			0.7696

Table 9: Italian BERT classification report for Tweets.

	precision	recall	f1-score
0	0.7290	0.9530	0.8261
1	0.8193	0.3757	0.5152
accuracy			0.7440
macro avg F1			0.6706

Table 10: Italian BERT classification report for News Headlines.

Both models had stable performances in the in-domain task, with the Italian BERT performing slightly better than the AlBERTo model, and both saw a significant decrease in performance for the cross-domain task, especially Italian BERT (macro F1: from 0.7696 to 0.6706). This is due to the recall dropping significantly for class 1 (Italian BERT: 0.9530 for class 0 and 0.3757 for class 1; AlBERTo: 0.9028 for class 0 and 0.4475 for class 1): once again, it proves difficult for the models to detect hate speech in news headlines.

It is worth mentioning that even after reducing the number of epochs to two, the validation loss still shows a slight upward trend for both models, as shown in Figures 8 and 9. It would be beneficial to delve deeper into

	precision	recall	f1-score
0	0.7805	0.7379	0.7586
1	0.7443	0.7862	0.7647
accuracy			0.7617
macro avg F1			0.7616

Table 11: AlBERTo classification report for Tweets.

	precision	recall	f1-score
0	0.7423	0.9028	0.8147
1	0.7232	0.4475	0.5529
accuracy			0.7380
macro avg F1			0.6838

Table 12: AlBERTo classification report for News headlines.

hyperparameter tuning and regularization techniques to see how much the performances of the models could improve.

4 DISCUSSION

As anticipated, the cross-domain task proved to be particularly challenging. All the tested models showed significantly lowered performance on the news dataset. Crucially, the recall of the positive class was unsatisfactory: this means that all models struggle with detecting hate speech in newspapers' headlines that are, indeed, hateful. Even the best performing model in this respect (the linear SVC trained on AlBERTo sentence embeddings) was able to detect hate speech in less than half of the cases.

Among all the features we developed, semantic information as modeled by sentence embeddings proved to be the most effective in the detection of hate speech. Even so, some more interpretable features also emerged, e.g. the count of bad words and the offensiveness score of documents, some unigrams like 'rom' and 'terrorismo', and the distribution of all uppercase words, which is commonly associated with expressions of rage.

Another notable, even if less obvious, stylistic feature, was the number of urls, which was strongly associated with the negative class. We hypothesize that urls occur
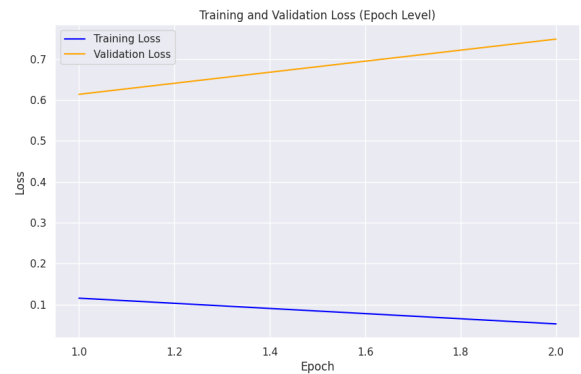

Figure 8: Italian BERT Training and Validation Loss trend (LR = 2e-5, Epochs = 2).

Figure 9: AlBERTo Training and Validation Loss trend (LR = 2e-5, Epochs = 2).

more frequently as links to articles in tweets reporting news from newspapers, and that such tweets are normally more objective and less prone to hateful rethoric. An in-depth investigation in the composition of a subset of our tweets dataset by Comandini and Patti [2] revealed that one third of the corpus was composed by references to newspaper's articles. More exactly, news accounted for more than half (51%) of the non hateful tweets, while the hateful messages were mostly comments from single users (88.29%). These finding corroborate our hypothesis.

It is also possible that the prominence of news in non hateful tweets in our training set hindered the detection of hateful headlines in the news dataset.

We were also able to confirm the observation by Comandini and Patti that nominal sentences, despite conveying a significant part of hate speech, are not in any way more frequent in hateful tweets. In fact, the distribution of nominal sentences was not a strong predictor of hate speech.

Finally, even though we ultimately decided against including them in any classifier, analyzing emojis offered valuable insights on emotional expressions in immigration-related topics. The most frequent emoji revealed a focus on negative emotions.

5 CONCLUSION

We created and tested several classifiers for the detection of hate speech in Italian tweets. We experimented different representations of the documents, developing several different features concerning formal properties of the documents (linguistic profiling features), presence of bad or offensive words, and semantic features, either explicit and relatively interpretable (n-grams) or implicit (embeddings). We also experimented with different models and fine-tuned two pre-trained Transformer-based models.

We took on the challenge of generalizing our findings across textual domains, testing our models trained on tweets against newspaper headlines and gained precious insight into the problems associated with this transfer.

We were able to clear both baselines provided by the organizers of the Evalita2020 task and obtained results comparable to those of the groups that originally participated in the task both in the in-domain task and in the cross-domain one.

REFERENCES

[1] Andrea Cimino, Lorenzo De Mattei, and Felice Dell'Orletta. "Multi-task Learning in Deep Neural Networks at EVALITA 2018". In: *EVALITA Evaluation of NLP and Speech Tools for Italian*. Ed. by Tommaso Caselli et al. Torino: Accademia University Press, 2018, pp. 86–95. ISBN: 978-88-319-7842-2 978-88-319-7869-9. DOI: 10.4000/books.aaccademia.4527. URL: https://books.openedition.org/aaccademia/4527 (visited on 06/04/2024).

[2] Gloria Comandini and Viviana Patti. "An Impossible Dialogue! Nominal Utterances and Populist Rhetoric in an Italian Twitter Corpus of Hate Speech against Immigrants". In: *Proceedings of the Third Workshop on Abusive Language Online*. Proceedings of the Third Workshop on Abusive Language Online. Florence, Italy: Association for Computational Linguistics, 2019, pp. 163–171. DOI: 10.18653/v1/W19-3518. URL: https://www.aclweb.org/anthology/W19-3518 (visited on 12/04/2024).

[3] European Commission. *The EU Code of conduct on countering illegal hate speech online.* URL: https://commission.europa.eu/strategy-and-policy/policies/justice-and-fundamental-rights/combatting-discrimination/racism-and-xenophobia/eu-code-conduct-countering-illegal-hate-speech-online_en (visited on 10/11/2024).

[4] Fabio Poletto et al. "Hate Speech Annotation: Analysis of an Italian Twitter Corpus". In: *Proceedings of the Fourth Italian Conference on Computational Linguistics CLiC-it 2017*. Ed. by Roberto Basili, Malvina Nissim, and Giorgio Satta. Torino: Accademia University Press, 2017, pp. 263–268. ISBN: 978-88-99982-94-2. DOI: 10.4000/books.aaccademia.2448. URL: https://books.openedition.org/aaccademia/2448 (visited on 12/12/2024).

[5] Marco Polignano et al. "AlBERTo: Italian BERT Language Understanding Model for NLP Challenging Tasks Based on Tweets". In: *Proceedings of the Sixth Italian Conference on Computational Linguistics (CLiC-it 2019)*. Vol. 2481. CEUR, 2019. URL: https://www.scopus.com/inward/record.uri?eid=2-s2.0-85074851349&partnerID=40&md5=7abed946e06f76b3825ae5e294ffac14.

[6] Peng Qi et al. "Stanza: A Python Natural Language Processing Toolkit for Many Human Languages". In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics:*

System Demonstrations. 2020. URL: https://nlp.
stanford.edu/pubs/qi2020stanza.pdf.

[7] Didier Reynders. *7th evaluation of the Code of Conduct*. Fact-sheet. European Commission, Nov. 2022. URL: https : / / commission . europa . eu / document/download/5dcc2a40- 785d- 43f0- b806- f065386395de _ en ? filename = Factsheet % 20 - %207th % 20monitoring % 20round % 20of % 20the % 20Code%20of%20Conduct.pdf.

[8] Manuela Sanguinetti et al. "HaSpeeDe 2@ EVALITA2020: Overview of the EVALITA 2020 Hate Speech Detection Task". In: *Proceedings of the 7th evaluation campaign of Natural Language Processing and Speech tools for Italian (EVALITA 2020)*. Ed. by Valerio Basile et al. Online: CEUR.org, 2020.

[9] Stefan Schweter. *Italian BERT and ELECTRA models*. Version 1.0.1. Nov. 2020. DOI: 10.5281/zenodo. 4263142. URL: https : / / doi . org / 10 . 5281 / zenodo.4263142.

[10] Alice Tontodimamma et al. "An Italian lexical resource for incivility detection in online discourses". In: *Quality & Quantity* 57.4 (Aug. 2023), pp. 3019–3037. ISSN: 0033-5177, 1573-7845. DOI: 10 . 1007 / s11135 - 022 - 01494 - 7. URL: https : //link . springer . com/10 . 1007 / s11135 - 022 - 01494-7 (visited on 11/29/2024).